

▼ Assignment 3 (Regularization) by vipin_2011MT22

Implement L2, L1, early stopping and drop out regularization for the previous assignment and write a report comparing the results (Accuracy, Precision, Recall and F1).

```
#importing the libraries
import pandas as pd
import numpy as np
from keras.models import Sequential
from keras.layers import Dense, Dropout
from keras.callbacks import EarlyStopping
from sklearn.model_selection import train_test_split
from keras.constraints import maxnorm
from tensorflow.keras import regularizers
from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score
```

```
#importing the dataset
dataframe = pd.read_csv("diabetes.csv")
dataframe.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
#spilting of dataset into feature and label
df_label = dataframe['Outcome']
df_features = dataframe.drop('Outcome', 1)
df_features.replace('?', -99999, inplace=True)
print(df_label.head())
print(df_features.head())
```

```
0    1
1    0
2    1
3    0
4    1
Name: Outcome, dtype: int64
   Pregnancies  Glucose  BloodPressure  ...    BMI  DiabetesPedigreeFunction  Age
0             6     148             72  ...   33.6              0.627         50
1             1      85             66  ...   26.6              0.351         31
2             8     183             64  ...   23.3              0.672         32
3             1      89             66  ...   28.1              0.167         21
4             0     137             40  ...   43.1              2.288         33

[5 rows x 8 columns]
```

```
#hot encoding the label dataset
label = []
for lab in df_label:
    if lab == 1:
        label.append([1, 0]) # class 1
    elif lab == 0:
        label.append([0, 1]) # class 0
```

```
data = np.array(df_features)
label = np.array(label)
print(data.shape,label.shape)
```

(768, 8) (768, 2)

```
#spilting dataset into testing and training
x_train, x_test, y_train, y_test = train_test_split(data, label, test_size=0.2, random_state=2)
x_train.shape
```

(614, 8)

▼ For l2

```
#building our Neural Network
model = Sequential()
model.add(Dense(500, input_dim=8, activation='sigmoid',kernel_regularizer=regularizers.l2(0.01), bias_regularizer=regularizers.l2(0.01),
               activity_regularizer=regularizers.l2(0.01)))
model.add(Dense(100, activation='sigmoid',kernel_regularizer=regularizers.l2(0.01), bias_regularizer=regularizers.l2(0.01),
               activity_regularizer=regularizers.l2(0.01)))
model.add(Dense(2, activation='softmax'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(x_train,y_train, epochs=1000, batch_size=70, validation_data=(x_test, y_test))
```

```
Epoch 1/1000
9/9 [=====] - 1s 62ms/step - loss: 4.9969 - accuracy: 0.5040 - val_loss: 4.1902 - val_accuracy: 0.5000
Epoch 2/1000
9/9 [=====] - 0s 9ms/step - loss: 4.0711 - accuracy: 0.6400 - val_loss: 3.4837 - val_accuracy: 0.6000
Epoch 3/1000
9/9 [=====] - 0s 7ms/step - loss: 3.3932 - accuracy: 0.6151 - val_loss: 2.8637 - val_accuracy: 0.6000
Epoch 4/1000
9/9 [=====] - 0s 7ms/step - loss: 2.7929 - accuracy: 0.6133 - val_loss: 2.3187 - val_accuracy: 0.6000
Epoch 5/1000
9/9 [=====] - 0s 7ms/step - loss: 2.2719 - accuracy: 0.6292 - val_loss: 1.8990 - val_accuracy: 0.6000
Epoch 6/1000
9/9 [=====] - 0s 7ms/step - loss: 1.8678 - accuracy: 0.6500 - val_loss: 1.5995 - val_accuracy: 0.6000
Epoch 7/1000
9/9 [=====] - 0s 9ms/step - loss: 1.5873 - accuracy: 0.6369 - val_loss: 1.4041 - val_accuracy: 0.6000
Epoch 8/1000
9/9 [=====] - 0s 8ms/step - loss: 1.4131 - accuracy: 0.6175 - val_loss: 1.2765 - val_accuracy: 0.6000
Epoch 9/1000
9/9 [=====] - 0s 8ms/step - loss: 1.3075 - accuracy: 0.6078 - val_loss: 1.1833 - val_accuracy: 0.6000
Epoch 10/1000
9/9 [=====] - 0s 10ms/step - loss: 1.2178 - accuracy: 0.6157 - val_loss: 1.1186 - val_accuracy: 0.6000
Epoch 11/1000
9/9 [=====] - 0s 9ms/step - loss: 1.1543 - accuracy: 0.6302 - val_loss: 1.0695 - val_accuracy: 0.6000
Epoch 12/1000
9/9 [=====] - 0s 8ms/step - loss: 1.0905 - accuracy: 0.6431 - val_loss: 1.0298 - val_accuracy: 0.6000
Epoch 13/1000
9/9 [=====] - 0s 8ms/step - loss: 1.0566 - accuracy: 0.6345 - val_loss: 1.0007 - val_accuracy: 0.6000
Epoch 14/1000
9/9 [=====] - 0s 8ms/step - loss: 1.0359 - accuracy: 0.6206 - val_loss: 0.9777 - val_accuracy: 0.6000
Epoch 15/1000
9/9 [=====] - 0s 8ms/step - loss: 1.0048 - accuracy: 0.6379 - val_loss: 0.9599 - val_accuracy: 0.6000
Epoch 16/1000
9/9 [=====] - 0s 9ms/step - loss: 0.9877 - accuracy: 0.6364 - val_loss: 0.9434 - val_accuracy: 0.6000
Epoch 17/1000
9/9 [=====] - 0s 8ms/step - loss: 0.9702 - accuracy: 0.6410 - val_loss: 0.9324 - val_accuracy: 0.6000
Epoch 18/1000
9/9 [=====] - 0s 9ms/step - loss: 0.9589 - accuracy: 0.6355 - val_loss: 0.9345 - val_accuracy: 0.6000
Epoch 19/1000
9/9 [=====] - 0s 8ms/step - loss: 0.9515 - accuracy: 0.6431 - val_loss: 0.9143 - val_accuracy: 0.6000
Epoch 20/1000
9/9 [=====] - 0s 8ms/step - loss: 0.9486 - accuracy: 0.6278 - val_loss: 0.9083 - val_accuracy: 0.6000
Epoch 21/1000
9/9 [=====] - 0s 7ms/step - loss: 0.9288 - accuracy: 0.6485 - val_loss: 0.9158 - val_accuracy: 0.6000
Epoch 22/1000
9/9 [=====] - 0s 7ms/step - loss: 0.9348 - accuracy: 0.6380 - val_loss: 0.8960 - val_accuracy: 0.6000
Epoch 23/1000
9/9 [=====] - 0s 7ms/step - loss: 0.9137 - accuracy: 0.6632 - val_loss: 0.9011 - val_accuracy: 0.6000
Epoch 24/1000
9/9 [=====] - 0s 7ms/step - loss: 0.9217 - accuracy: 0.6200 - val_loss: 0.8896 - val_accuracy: 0.6000
Epoch 25/1000
9/9 [=====] - 0s 7ms/step - loss: 0.9012 - accuracy: 0.6578 - val_loss: 0.8955 - val_accuracy: 0.6000
Epoch 26/1000
9/9 [=====] - 0s 7ms/step - loss: 0.9006 - accuracy: 0.6633 - val_loss: 0.8850 - val_accuracy: 0.6000
Epoch 27/1000
9/9 [=====] - 0s 8ms/step - loss: 0.9010 - accuracy: 0.6418 - val_loss: 0.8826 - val_accuracy: 0.6000
Epoch 28/1000
9/9 [=====] - 0s 7ms/step - loss: 0.9141 - accuracy: 0.6244 - val_loss: 0.8802 - val_accuracy: 0.6000
Epoch 29/1000
9/9 [=====] - 0s 8ms/step - loss: 0.9019 - accuracy: 0.6428 - val_loss: 0.8752 - val_accuracy: 0.6000
Epoch 30/1000
9/9 [=====] - 0s 8ms/step - loss: 0.9019 - accuracy: 0.6428 - val_loss: 0.8752 - val_accuracy: 0.6000
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 500)	4500
dense_1 (Dense)	(None, 100)	50100
dense_2 (Dense)	(None, 2)	202
=====		
Total params: 54,802		
Trainable params: 54,802		
Non-trainable params: 0		

```
#predictating whether diabetic or not
Y_pred = model.predict(x_test)
Y_pred =(Y_pred>0.5)

#Calculating accuracy, f1 score, precision, recall

accuracyL2 = accuracy_score(y_test, Y_pred)
scoreL2 = f1_score(y_test, Y_pred, average='weighted')
precisionL2 = precision_score(y_test, Y_pred, labels=[0,1], average='weighted')
recallL2 = recall_score(y_test, Y_pred, average='weighted')

print('accuracy: ',accuracyL2)
print('F1-score: ',scoreL2)
print('Precision: ', precisionL2)
print("recall: ", recallL2)

accuracy:  0.7662337662337663
F1-score:  0.7589679907832293
Precision:  0.7563457563457563
recall:  0.7662337662337663
```

▼ For l1

```
#building our Neural Network
model = Sequential()
model.add(Dense(500, input_dim=8, activation='sigmoid',kernel_regularizer=regularizers.l1(0.0001), bias_regularizer=regularizers.l1(0.0001), activity_regularizer=regularizers.l1(0.0001)))
model.add(Dense(100, activation='sigmoid',kernel_regularizer=regularizers.l1(0.0001), bias_regularizer=regularizers.l1(0.0001), activity_regularizer=regularizers.l1(0.0001)))
model.add(Dense(2, activation='softmax'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(x_train,y_train, epochs=1000, batch_size=70, validation_data=(x_test, y_test))

Epoch 1/1000
9/9 [=====] - 1s 29ms/step - loss: 1.0600 - accuracy: 0.5612 - val_loss: 0.9176 - val_accuracy: 0.5612
Epoch 2/1000
9/9 [=====] - 0s 8ms/step - loss: 0.9787 - accuracy: 0.6179 - val_loss: 0.8942 - val_accuracy: 0.6179
Epoch 3/1000
9/9 [=====] - 0s 8ms/step - loss: 0.9042 - accuracy: 0.6539 - val_loss: 0.8453 - val_accuracy: 0.6539
Epoch 4/1000
9/9 [=====] - 0s 8ms/step - loss: 0.8635 - accuracy: 0.6711 - val_loss: 0.8345 - val_accuracy: 0.6711
Epoch 5/1000
9/9 [=====] - 0s 8ms/step - loss: 0.8565 - accuracy: 0.6731 - val_loss: 0.8156 - val_accuracy: 0.6731
Epoch 6/1000
9/9 [=====] - 0s 7ms/step - loss: 0.8343 - accuracy: 0.6860 - val_loss: 0.7975 - val_accuracy: 0.6860
Epoch 7/1000
9/9 [=====] - 0s 8ms/step - loss: 0.8215 - accuracy: 0.6943 - val_loss: 0.7809 - val_accuracy: 0.6943
Epoch 8/1000
9/9 [=====] - 0s 9ms/step - loss: 0.8143 - accuracy: 0.7033 - val_loss: 0.7662 - val_accuracy: 0.7033
Epoch 9/1000
9/9 [=====] - 0s 9ms/step - loss: 0.8068 - accuracy: 0.7045 - val_loss: 0.7644 - val_accuracy: 0.7045
Epoch 10/1000
9/9 [=====] - 0s 7ms/step - loss: 0.7747 - accuracy: 0.7137 - val_loss: 0.7473 - val_accuracy: 0.7137
Epoch 11/1000
9/9 [=====] - 0s 8ms/step - loss: 0.7582 - accuracy: 0.7369 - val_loss: 0.7686 - val_accuracy: 0.7369
Epoch 12/1000
9/9 [=====] - 0s 8ms/step - loss: 0.7373 - accuracy: 0.7479 - val_loss: 0.7323 - val_accuracy: 0.7479
Epoch 13/1000
9/9 [=====] - 0s 8ms/step - loss: 0.7367 - accuracy: 0.7280 - val_loss: 0.7340 - val_accuracy: 0.7280
Epoch 14/1000
9/9 [=====] - 0s 8ms/step - loss: 0.7391 - accuracy: 0.7368 - val_loss: 0.7247 - val_accuracy: 0.7368
Epoch 15/1000
9/9 [=====] - 0s 8ms/step - loss: 0.7191 - accuracy: 0.7461 - val_loss: 0.7163 - val_accuracy: 0.7461
Epoch 16/1000
9/9 [=====] - 0s 8ms/step - loss: 0.7182 - accuracy: 0.7630 - val_loss: 0.7379 - val_accuracy: 0.7630
Epoch 17/1000
9/9 [=====] - 0s 8ms/step - loss: 0.7052 - accuracy: 0.7467 - val_loss: 0.7176 - val_accuracy: 0.7467
Epoch 18/1000
9/9 [=====] - 0s 10ms/step - loss: 0.6985 - accuracy: 0.7490 - val_loss: 0.6953 - val_accuracy: 0.7490
Epoch 19/1000
9/9 [=====] - 0s 9ms/step - loss: 0.6904 - accuracy: 0.7667 - val_loss: 0.6999 - val_accuracy: 0.7667
Epoch 20/1000
9/9 [=====] - 0s 8ms/step - loss: 0.6912 - accuracy: 0.7518 - val_loss: 0.7041 - val_accuracy: 0.7518
Epoch 21/1000
9/9 [=====] - 0s 8ms/step - loss: 0.6894 - accuracy: 0.7462 - val_loss: 0.6843 - val_accuracy: 0.7462
Epoch 22/1000
9/9 [=====] - 0s 8ms/step - loss: 0.6813 - accuracy: 0.7644 - val_loss: 0.6943 - val_accuracy: 0.7644
Epoch 23/1000
9/9 [=====] - 0s 7ms/step - loss: 0.6714 - accuracy: 0.7593 - val_loss: 0.7148 - val_accuracy: 0.7593
Epoch 24/1000
```

4/11/2021DL_Assignment3.ipynb - Colaboratory

9/9 [=====] - 0s 7ms/step - loss: 0.7014 - accuracy: 0.7237 - val_loss: 0.6805 - val_accuracy: 0.7434
Epoch 25/1000
9/9 [=====] - 0s 8ms/step - loss: 0.6752 - accuracy: 0.7502 - val_loss: 0.6648 - val_accuracy: 0.7605
Epoch 26/1000
9/9 [=====] - 0s 7ms/step - loss: 0.6654 - accuracy: 0.7594 - val_loss: 0.7324 - val_accuracy: 0.7434
Epoch 27/1000
9/9 [=====] - 0s 7ms/step - loss: 0.6574 - accuracy: 0.7617 - val_loss: 0.6671 - val_accuracy: 0.7605
Epoch 28/1000
9/9 [=====] - 0s 7ms/step - loss: 0.6464 - accuracy: 0.7605 - val_loss: 0.6639 - val_accuracy: 0.7605
Epoch 29/1000
9/9 [=====] - 0s 7ms/step - loss: 0.6682 - accuracy: 0.7434 - val_loss: 0.6688 - val_accuracy: 0.7605

model.summary()

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 500)	4500
dense_4 (Dense)	(None, 100)	50100
dense_5 (Dense)	(None, 2)	202

Total params: 54,802
Trainable params: 54,802
Non-trainable params: 0

#predictating whether diabetic or not
Y_pred = model.predict(x_test)
Y_pred =(Y_pred>0.5)

#Calculating accuracy, f1 score, precision, recall

accuracyL1 = accuracy_score(y_test, Y_pred)
scoreL1 = f1_score(y_test, Y_pred, average='weighted')
precisionL1 = precision_score(y_test, Y_pred, labels=[0,1], average='weighted')
recallL1 = recall_score(y_test, Y_pred, average='weighted')

print('accuracy: ',accuracyL1)
print('F1-score: ',scoreL1)
print('Precision: ', precisionL1)
print("recall: ", recallL1)

accuracy: 0.7077922077922078
F1-score: 0.708724455268234
Precision: 0.7097162097162096
recall: 0.7077922077922078

▼ For Early Stopping

#building our Neural Network
model = Sequential()
model.add(Dense(500, input_dim=8, activation='sigmoid'))
model.add(Dense(100, activation='sigmoid'))
model.add(Dense(2, activation='softmax'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

es = EarlyStopping(monitor='val_loss',patience=300)
model.fit(x_train,y_train, epochs=1000, batch_size=70, validation_data=(x_test,y_test),callbacks=[es])

9/9 [=====] - 0s 9ms/step - loss: 0.2341 - accuracy: 0.9007 - val_loss: 0.7982 - val_accuracy: 0.8915
Epoch 346/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2364 - accuracy: 0.8915 - val_loss: 0.7903 - val_accuracy: 0.8915
Epoch 347/1000
9/9 [=====] - 0s 9ms/step - loss: 0.2240 - accuracy: 0.8957 - val_loss: 0.8117 - val_accuracy: 0.8957
Epoch 348/1000
9/9 [=====] - 0s 7ms/step - loss: 0.2397 - accuracy: 0.8932 - val_loss: 0.8049 - val_accuracy: 0.8932
Epoch 349/1000
9/9 [=====] - 0s 7ms/step - loss: 0.2334 - accuracy: 0.9035 - val_loss: 0.9362 - val_accuracy: 0.9035
Epoch 350/1000
9/9 [=====] - 0s 9ms/step - loss: 0.2353 - accuracy: 0.9094 - val_loss: 0.8247 - val_accuracy: 0.9094
Epoch 351/1000
9/9 [=====] - 0s 7ms/step - loss: 0.2275 - accuracy: 0.9045 - val_loss: 0.8382 - val_accuracy: 0.9045
Epoch 352/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2331 - accuracy: 0.9018 - val_loss: 0.7882 - val_accuracy: 0.9018
Epoch 353/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2006 - accuracy: 0.9240 - val_loss: 0.7753 - val_accuracy: 0.9240
Epoch 354/1000
9/9 [=====] - 0s 7ms/step - loss: 0.2167 - accuracy: 0.9035 - val_loss: 0.8099 - val_accuracy: 0.9035
Epoch 355/1000

4/11/2021DL_Assignment3.ipynb - Colaboratory

```
-- -- -- -- --
9/9 [=====] - 0s 7ms/step - loss: 0.2490 - accuracy: 0.8927 - val_loss: 0.8004 - val_accuracy: 0.8927
Epoch 356/1000
9/9 [=====] - 0s 7ms/step - loss: 0.2357 - accuracy: 0.9090 - val_loss: 0.8150 - val_accuracy: 0.9090
Epoch 357/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2309 - accuracy: 0.9043 - val_loss: 0.8188 - val_accuracy: 0.9043
Epoch 358/1000
9/9 [=====] - 0s 7ms/step - loss: 0.1891 - accuracy: 0.9276 - val_loss: 0.8116 - val_accuracy: 0.9276
Epoch 359/1000
9/9 [=====] - 0s 9ms/step - loss: 0.2212 - accuracy: 0.9106 - val_loss: 0.8707 - val_accuracy: 0.9106
Epoch 360/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2220 - accuracy: 0.9107 - val_loss: 0.9720 - val_accuracy: 0.9107
Epoch 361/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2702 - accuracy: 0.8761 - val_loss: 0.9762 - val_accuracy: 0.8761
Epoch 362/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2567 - accuracy: 0.8849 - val_loss: 1.0803 - val_accuracy: 0.8849
Epoch 363/1000
9/9 [=====] - 0s 7ms/step - loss: 0.3259 - accuracy: 0.8672 - val_loss: 0.8595 - val_accuracy: 0.8672
Epoch 364/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2495 - accuracy: 0.8837 - val_loss: 0.8244 - val_accuracy: 0.8837
Epoch 365/1000
9/9 [=====] - 0s 7ms/step - loss: 0.2265 - accuracy: 0.9077 - val_loss: 0.8118 - val_accuracy: 0.9077
Epoch 366/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2268 - accuracy: 0.8958 - val_loss: 0.7841 - val_accuracy: 0.8958
Epoch 367/1000
9/9 [=====] - 0s 9ms/step - loss: 0.1996 - accuracy: 0.9216 - val_loss: 0.8105 - val_accuracy: 0.9216
Epoch 368/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2114 - accuracy: 0.9119 - val_loss: 0.8484 - val_accuracy: 0.9119
Epoch 369/1000
9/9 [=====] - 0s 8ms/step - loss: 0.1854 - accuracy: 0.9367 - val_loss: 0.8963 - val_accuracy: 0.9367
Epoch 370/1000
9/9 [=====] - 0s 7ms/step - loss: 0.2048 - accuracy: 0.9291 - val_loss: 0.8741 - val_accuracy: 0.9291
Epoch 371/1000
9/9 [=====] - 0s 7ms/step - loss: 0.2100 - accuracy: 0.9130 - val_loss: 0.8298 - val_accuracy: 0.9130
Epoch 372/1000
9/9 [=====] - 0s 7ms/step - loss: 0.2459 - accuracy: 0.8966 - val_loss: 0.9344 - val_accuracy: 0.8966
Epoch 373/1000
9/9 [=====] - 0s 8ms/step - loss: 0.2932 - accuracy: 0.8613 - val_loss: 0.9151 - val_accuracy: 0.8613
<tensorflow.python.keras.callbacks.History at 0x7f9024080550>
```

model.summary()

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
dense_21 (Dense)	(None, 500)	4500

dense_22 (Dense)	(None, 100)	50100

dense_23 (Dense)	(None, 2)	202
=====		
Total params: 54,802		
Trainable params: 54,802		
Non-trainable params: 0		

#predictating whether diabetic or not
Y_pred = model.predict(x_test)
Y_pred =(Y_pred>0.5)

#Calculating accuracy, f1 score, precision, recall

accuracyES = accuracy_score(y_test, Y_pred)
scoreES = f1_score(y_test, Y_pred, average='weighted')
precisionES = precision_score(y_test, Y_pred, labels=[0,1], average='weighted')
recallES = recall_score(y_test, Y_pred, average='weighted')

print('accuracy: ',accuracyES)
print('F1-score: ',scoreES)
print('Precision: ', precisionES)
print("recall: ", recallES)

accuracy: 0.7402597402597403
F1-score: 0.724173285879275
Precision: 0.7220922272063107
recall: 0.7402597402597403

▼ For Dropout

```
model = Sequential()  
model.add(Dense(500, input_dim=8, activation='sigmoid',kernel_constraint=maxnorm(3)))  
model.add(Dropout(0.3))  
  
model.add(Dense(100, activation='sigmoid',kernel_constraint=maxnorm(3)))  
model.add(Dropout(0.3))  
  
model.add(Dense(2, activation='softmax'))  
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])  
model.fit(x_train,y_train, epochs=1000, batch_size=70, validation_data=(x_test, y_test))
```

Epoch 1/1000
9/9 [=====] - 1s 23ms/step - loss: 0.6998 - accuracy: 0.5995 - val_loss: 0.5902 - val_accuracy: 0.5995
Epoch 2/1000
9/9 [=====] - 0s 7ms/step - loss: 0.6881 - accuracy: 0.6176 - val_loss: 0.5792 - val_accuracy: 0.6176
Epoch 3/1000
9/9 [=====] - 0s 8ms/step - loss: 0.6518 - accuracy: 0.6068 - val_loss: 0.5746 - val_accuracy: 0.6068
Epoch 4/1000
9/9 [=====] - 0s 8ms/step - loss: 0.6517 - accuracy: 0.6242 - val_loss: 0.5681 - val_accuracy: 0.6242
Epoch 5/1000
9/9 [=====] - 0s 9ms/step - loss: 0.6012 - accuracy: 0.6751 - val_loss: 0.5571 - val_accuracy: 0.6751
Epoch 6/1000
9/9 [=====] - 0s 9ms/step - loss: 0.6231 - accuracy: 0.6482 - val_loss: 0.5566 - val_accuracy: 0.6482
Epoch 7/1000
9/9 [=====] - 0s 7ms/step - loss: 0.6039 - accuracy: 0.6847 - val_loss: 0.5504 - val_accuracy: 0.6847
Epoch 8/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5850 - accuracy: 0.6991 - val_loss: 0.5463 - val_accuracy: 0.6991
Epoch 9/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5824 - accuracy: 0.6975 - val_loss: 0.5572 - val_accuracy: 0.6975
Epoch 10/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5762 - accuracy: 0.7182 - val_loss: 0.5374 - val_accuracy: 0.7182
Epoch 11/1000
9/9 [=====] - 0s 8ms/step - loss: 0.5976 - accuracy: 0.6782 - val_loss: 0.5479 - val_accuracy: 0.6782
Epoch 12/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5646 - accuracy: 0.6917 - val_loss: 0.5371 - val_accuracy: 0.6917
Epoch 13/1000
9/9 [=====] - 0s 8ms/step - loss: 0.5878 - accuracy: 0.6747 - val_loss: 0.5353 - val_accuracy: 0.6747
Epoch 14/1000
9/9 [=====] - 0s 8ms/step - loss: 0.5468 - accuracy: 0.7158 - val_loss: 0.5537 - val_accuracy: 0.7158
Epoch 15/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5689 - accuracy: 0.7125 - val_loss: 0.5325 - val_accuracy: 0.7125
Epoch 16/1000
9/9 [=====] - 0s 10ms/step - loss: 0.5440 - accuracy: 0.7412 - val_loss: 0.5311 - val_accuracy: 0.7412
Epoch 17/1000
9/9 [=====] - 0s 8ms/step - loss: 0.5440 - accuracy: 0.7321 - val_loss: 0.5186 - val_accuracy: 0.7321
Epoch 18/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5248 - accuracy: 0.7637 - val_loss: 0.5213 - val_accuracy: 0.7637
Epoch 19/1000
9/9 [=====] - 0s 8ms/step - loss: 0.5450 - accuracy: 0.7438 - val_loss: 0.5175 - val_accuracy: 0.7438
Epoch 20/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5159 - accuracy: 0.7626 - val_loss: 0.5364 - val_accuracy: 0.7626
Epoch 21/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5564 - accuracy: 0.7509 - val_loss: 0.5197 - val_accuracy: 0.7509
Epoch 22/1000
9/9 [=====] - 0s 8ms/step - loss: 0.5570 - accuracy: 0.7269 - val_loss: 0.5166 - val_accuracy: 0.7269
Epoch 23/1000
9/9 [=====] - 0s 8ms/step - loss: 0.5191 - accuracy: 0.7437 - val_loss: 0.5577 - val_accuracy: 0.7437
Epoch 24/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5401 - accuracy: 0.7234 - val_loss: 0.5150 - val_accuracy: 0.7234
Epoch 25/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5066 - accuracy: 0.7769 - val_loss: 0.5169 - val_accuracy: 0.7769
Epoch 26/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5141 - accuracy: 0.7661 - val_loss: 0.5251 - val_accuracy: 0.7661
Epoch 27/1000
9/9 [=====] - 0s 8ms/step - loss: 0.5287 - accuracy: 0.7308 - val_loss: 0.5107 - val_accuracy: 0.7308
Epoch 28/1000
9/9 [=====] - 0s 8ms/step - loss: 0.4971 - accuracy: 0.7694 - val_loss: 0.5280 - val_accuracy: 0.7694
Epoch 29/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5251 - accuracy: 0.7344 - val_loss: 0.5218 - val_accuracy: 0.7344
Epoch 30/1000
9/9 [=====] - 0s 7ms/step - loss: 0.5251 - accuracy: 0.7344 - val_loss: 0.5218 - val_accuracy: 0.7344

```
model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 500)	4500
dropout (Dropout)	(None, 500)	0
dense_10 (Dense)	(None, 100)	50100
dropout_1 (Dropout)	(None, 100)	0
dense_11 (Dense)	(None, 2)	202
Total params: 54,802		

Trainable params: 54,802
Non-trainable params: 0

```
#predictating whether diabetic or not
Y_pred = model.predict(x_test)
Y_pred =(Y_pred>0.5)

#Calculating accuracy, f1 score, precision, recall

accuracyD0 = accuracy_score(y_test, Y_pred)
scoreD0 = f1_score(y_test, Y_pred, average='weighted')
precisionD0 = precision_score(y_test, Y_pred, labels=[0,1], average='weighted')
recallD0 = recall_score(y_test, Y_pred, average='weighted')

print('accuracy: ',accuracyD0)
print('F1-score: ',scoreD0)
print('Precision: ', precisionD0)
print("recall: ", recallD0)

accuracy:  0.7077922077922078
F1-score:  0.7178683385579938
Precision:  0.7409505388228792
recall:  0.7077922077922078
```

▼ Report

```
print('For L2')
print('accuracy: ',accuracyL2)
print('F1-score: ',scoreL2)
print('Precision: ', precisionL2)
print("recall: ", recallL2, '\n')

print('For L1')
print('accuracy: ',accuracyL1)
print('F1-score: ',scoreL1)
print('Precision: ', precisionL1)
print("recall: ", recallL1, '\n')

print('For Early Stopping')
print('accuracy: ',accuracyES)
print('F1-score: ',scoreES)
print('Precision: ', precisionES)
print("recall: ", recallES, '\n')

print('For Dropout')
print('accuracy: ',accuracyD0)
print('F1-score: ',scoreD0)
print('Precision: ', precisionD0)
print("recall: ", recallD0, '\n')

For L2
accuracy:  0.7662337662337663
F1-score:  0.7589679907832293
Precision:  0.7563457563457563
recall:  0.7662337662337663

For L1
accuracy:  0.7077922077922078
F1-score:  0.708724455268234
Precision:  0.7097162097162096
recall:  0.7077922077922078

For Early Stopping
accuracy:  0.7402597402597403
F1-score:  0.724173285879275
Precision:  0.7220922272063107
recall:  0.7402597402597403

For Dropout
accuracy:  0.7077922077922078
F1-score:  0.7178683385579938
Precision:  0.7409505388228792
recall:  0.7077922077922078
```

L2 regularization model is giving maximum accuracy followed by EarlyStopping. Dropout and L1 regularization model is giving same accuracy.

F1-score of L2 regularization model is maximum followed by Dropout regularization model, EarlyStopping regularization model and L1 regularization model.

Precision of L2 regularization model is maximum followed by Dropout regularization model, L1 regularization model and EarlyStopping regularization model.

L2 regularization model is giving maximum recall followed by EarlyStopping regularization model, Dropout regularization model and L1 regularization model

for early stopping, the patience parameter plays a huge role, for patience = 300, model stopped at 373 epochs with accuracy of around 74% on test data. This gives a hint that epochs = 1000 is bound to give an overfitted model

comparing between L1 and L2 when for the same models we used a parameter of 0.01 for both the regularisations, L2 had better results and L1 seemed overfitted

overall the metrics are comparable for all models however the L2 model seems to perform slightly better than other models

